

**Learn to Schedule (LEASCH): A Deep reinforcement learning approach
for radio resource scheduling in the 5G MAC layer
&
Cellular Network Traffic Scheduling with Deep Reinforcement Learning**

Presented by: Salman Mohebi
ITN WindMill Study Group Meeting
June 10, 2020

Learn to Schedule (LEASCH): A Deep reinforcement learning approach for radio resource scheduling in the 5G MAC layer.

F. AL-Tam*, N. Correia*, J. Rodriguez †

*Centro de Electrónica, Optoelectrónica e Telecomunicações (CEOT) - Universidade do Algarve,

†Instituto de Telecomunicações, Universidade de Aveiro

Main contributions

- DRL model for Radio resource scheduling (RRS) in the 5G MAC layer:
 - **learn-rather-than-design (LRTD) approach**
- Pipeline for developing/training DRL agents
- Analysis the proposed model vs. baseline algorithms in different network settings

Reinforcement learning

- Markov decision process (MDP): $(\mathcal{S}, \mathcal{A}, \mathbf{P}, r, \gamma)$.

- Goal: maximize $G_t = \mathbb{E} \left[\sum_{k=0}^{\infty} \gamma^k r(s_{t+k}, a_{t+k}) \mid s_0 = s_t \right]$

- state-value, function $V(s) = \mathbb{E}[G_t \mid s_t = s]$
- action-value function $Q(s, a) = \mathbb{E}[G_t \mid s_t = s, a_t = a]$
- advantage function $A(s, a) = Q(s, a) - V(s)$.

$$V(s) = \sum_{a \in \mathcal{A}} \pi(a|s) Q(s, a)$$

Reinforcement learning

- Bellman expectation function

$$Q(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}(a) V(s')$$

- Bellman optimality equation

$$Q^*(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p_{ss'}(a) \max_{a'} Q^*(s', a')$$

- Optimal policy

$$\pi^*(s) = \arg \max_{a \in \mathcal{A}} Q^*(s, a), \quad \forall s \in \mathcal{S}$$

- Problem: Transition probability is not known
 - Q- learning algorithm

$$Q(s_t, a_t) = r(s_t, a_t) + \gamma \max_{a_{t+1}} Q(s_{t+1}, a_{t+1})$$

$$\text{loss} = (Q(s_t, a_t; \theta) - Q^{\text{target}})^2$$

- Approximate Q

$$Q^{\text{target}} = r(s_t, a_t) + \gamma \max_{a_{t+1}} Q(s_{t+1}, a_{t+1}; \theta)$$

- Deep Q networks (DQN)
- Mean-squared Bellman error (MSBE)

$$\theta_t = \theta_t + \frac{\alpha}{M} (Q(s, a; \theta_t) - Q^{\text{target}}(\theta_t)) \nabla_{\theta_t} Q(s, a; \theta_t)$$

Reinforcement learning


- Stabilize the results:
 - Two identical neural networks are: **target network** and **online network**

$$Q^{\text{target}} = r(s_t, a_t) + \gamma \max_{a_{t+1}} Q(s_{t+1}, a_{t+1}; \hat{\theta})$$

- Experience replay memory $\hat{\theta} = \beta \theta + (1 - \beta) \hat{\theta}$
- Double DQN (DDQN)
$$Q^{\text{target}} = r_{t+1}(s, a) + \gamma Q(s_{t+1}, \arg \max_a Q(s_{t+1}, a, \theta); \hat{\theta})$$

LEASCH's design-1

- Scheduler runs at the gNB at every slot
 - Grant the available resource block groups (RBGs) between UEs
 - Filling the resource grid
- Goal: jointly optimize the throughput and fairness
- State space:
 - Eligibility $g_u = \begin{cases} 1, & \text{if } u \text{ is eligible} \\ 0, & \text{Otherwise} \end{cases}, \forall u \in \mathcal{U}$
 - Data Rate (d)
 - Fairness (f) $f_u = \begin{cases} \max(f_u - 1, 0), & \text{if } u \text{ is selected} \\ f_u + 1, & \text{if } u\text{'s buffer is not empty} \end{cases}, \forall u \in \mathcal{U}$
- Compact State: $\hat{d} = d \circ g$ $s = [\hat{d} \quad f]^\top$
- Normalize \hat{d} and f to the range $[0, 1]$



LEASCH's design-2

- Action space: select one of the UEs in the system
- Reward:
 - Encourage the agent to transmit at the RBGs with the highest MCS (throughput)
 - Not to compromise the resource sharing between the users (fairness)

$$r(s, u; K) = \begin{cases} -K, & \text{if } u \text{ is none-eligible} \\ \hat{d}_u \times \frac{\min_u f_u}{\max_u f_u}, & \text{otherwise} \end{cases} \quad (21)$$

- K is a threshold to penalize the scheduling an inactive UE

LEASCH training and deployment

Algorithm 1 - Training phase of LEASCH.

```

1: // input:  $\ell_{\text{episode}}, K, M, T \in, \delta_\epsilon, \min_\epsilon, \theta, \hat{\theta}, \mathcal{R}$ .
2: // output: updated  $\{\theta, \hat{\theta}, \mathcal{R}\}$ .
3: initialize  $s$  randomly according to the ranges of  $\hat{d}$  and  $f$ 
4: for  $i = 1 : \ell_{\text{episode}}$  do
5:   forward  $s$  to the on-line Q neural network and get the
     selected UE,  $u$ , via  $\epsilon$ -greedy as:
     
$$u = \arg \max_{a \in \mathcal{A}} Q(s, a; \theta)$$

6:  $\rightarrow$  anneal  $\epsilon$  as:  $\max\{\epsilon - \delta_\epsilon, \min_\epsilon\}$ 
7:   calculate the reward  $r(s, u; K)$  using (21)
8:   calculate new state  $s'$  using the equations  $s = [\hat{d} \ f]^\top$ 
9:   add the tuple  $(s, u, r, s')$  to the experience replay  $\mathcal{R}$ 
10:  sample  $M$  mini-batches from  $\mathcal{R}$  and train the on-line
     Q neural network with  $\theta$  using (14) and (17)
11:  $\rightarrow$  update the target critic Q neural network (with  $\hat{\theta}$ ) using
      $\theta$  every  $T$  steps via smoothing (16).
12:   $s \leftarrow s'$ 
13: end for
14: return  $\{\theta, \hat{\theta}, \mathcal{R}\}$ 

```

Algorithm 2 - Deployment phase of LEASCH in 5G.

```

1: // input: trained LEASCH.
2: for each time slot do
3:   for each RBG do
4:     calculate the set of eligible UEs  $\hat{\mathcal{U}}$ 
5:     if  $\hat{\mathcal{U}} \neq \emptyset$  then
6:       calculate state  $s$ 
7:       forward  $s$  to LEASCH
8:       calculate the action  $u$  as:
         
$$u = \arg \max_{a \in \mathcal{A}} Q(s, a; \theta)$$

9:       if  $u \in \hat{\mathcal{U}}$  then
10:         schedule  $u$  for the current RBG
11:       end if
12:     end if
13:   collect statistics from the simulator
14: end for
15: end for

```

Simulation parameters

| Parameter | Value | Description |
|-------------------------|-----------|-------------------------------|
| α | $1e^{-4}$ | DNN learning rate |
| Optimizer | Adam | |
| Gradient threshold | 1 | |
| ϵ | 0.99 | ϵ -greedy parameter |
| \min_{ϵ} | 0.01 | Min. allowed ϵ |
| δ_{ϵ} | $1e^{-4}$ | ϵ decaying factor |
| $ \mathcal{R} $ | $1e^6$ | Experience replay memory size |
| M | 64 | Mini-batch size |
| T | 20 | Smoothing frequency |
| β | $1e^{-3}$ | Smoothing threshold |
| ℓ_{episode} | 150 RBG | Episode length |
| No. of episodes | 500 | Training episodes |

| Parameter | Value |
|------------------------|--|
| Radio access tech. | 3GPP 5G NR |
| Test time | 250 frames |
| Simulation runs | 100 runs with different deployment scenarios |
| Numerology index μ | $\{0, 1, 2\}$ |
| Bandwidth | $\{5\text{MHz}, 10\text{MHz}, 20\text{MHz}\}$ |
| UEs | 4 |
| SCS | $\{15\text{kHz}, 30\text{kHz}, 60\text{kHz}\}$ |
| No. of RBs | $\{25, 24, 24\}$ see [2] |
| Scheduling period | 1 RGB |
| RBG size | 2 RBs according to configuration 1 in [3] |
| Total tested RBGs | $250 \times 100 \times \{130, 240, 480\}$ RBGs |
| Channel development | Randomly changes each $\frac{1}{4}$ second |
| HARQ | True |

Simulation setup

- Q neural networks are DNNs:
 - Two fully connected hidden layers of 128 neurons each
 - Relu activation functions
 - Input layer size: $2 \times |U|$
 - Output layer size: $|U|$
- Matlab 2019b:
 - Linux with i7 2.6GHz,
 - 32GB RAM,
 - GPU Nvidia RTX 2080Ti with 11 GB
- Key performance indicators:
 - Throughput: achievable data rate in the cell
 - Goodput: Delivered data rate at receiver
 - Fairness: Jain's fairness index

Results




Figure 2: LEASCH learning curve for 2000 episodes.




Figure 3: KPIs for 250 frames of 15kHz SCS under 5MHz BW for 100 runs.




Figure 4: KPIs for 250 frames of 30kHz SCS under 10MHz BW for 100 runs.




Figure 5: KPIs for 250 frames of 60kHz SCS under 20MHz BW for 100 runs.

Results

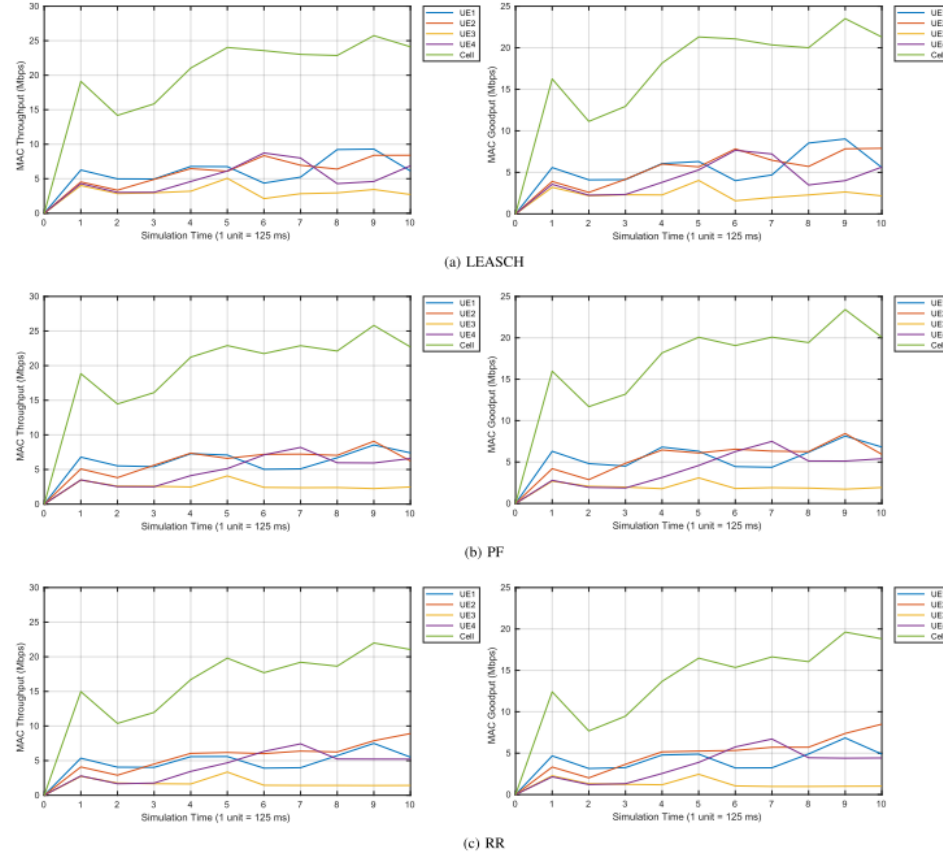


Figure 6: A random testing run for the 10MHz BW and 30kHz SCS setting. Left column: throughput; right column: goodput.

Cellular Network Traffic Scheduling with Deep Reinforcement Learning

Sandeep Chinchali¹, **Pan Hu**², **Tianshu Chu**³, **Manu Sharma**³, **Manu Bansal**³, **Rakesh Misra**³
Marco Pavone⁴ and **Sachin Katti**^{1,2}

¹ Department of Computer Science, Stanford University

² Department of Electrical Engineering, Stanford University

³ Uhana, Inc.

⁴ Department of Aeronautics and Astronautics, Stanford University


{csandeep, panhu, pavone, skatti}@stanford.edu, {tchu, manusharma, manub, rakesh}@uhana.io

Introduction

- Focus on mobile networks,
 - Increasing request for new class of applications, driven by IoT
- High Volume Flexible Time (HVFT) applications
 - Software and data updates to mobile IoT devices:
 - Updating maps for self-driving cars or delivery drones
 - Large transfer of IoT sensor data to the cloud:
 - Energy usage measurements from a smart grid,
 - Prefetched ultra-high quality and bitrate video

Challenges

- Time-variant and non-Markovian network dynamics
- Networks exhibit non-stationary dynamics
 - ranging from short-term, minute-scale variation to daily commute patterns
- Data set: 4 weeks of data from 10 diverse cells in Melbourne, Australia
- Incorporate **past measurements** and **historical** commute patterns into our state representation to recast the problem as MDP to leverage RL methods:
 - Temporal feature mapping function



Data-driven Network Model-1

- Discrete-time, continuous state and action space MDP,

- **State space:**

$$S_t = [C_t, N_t, E_t]$$

- Cell Congestion (C): The effective number of users in the cell
- Average Cell Efficiency (E): average cell quality
 - Total cell bandwidth, type of cellular technology deployed, distance of users from the cell tower, and ...
- Number of connections (N)

$$s_t = [S_t, \phi(S_0, \dots, S_t, t, T)]$$

- ϕ is a **temporal feature mapping** function

- feed-forward neural networks,
- Long Short Term Memory (LSTM) networks,
- Autoregressive Integrated Moving Average (ARIMA) techniques

Data-driven Network Model-2

- Action: $a_t \in [0, 1]$
 - rate at which HVFT traffic can be served on top of conventional traffic

- Reward:

$$R(s_t, a_t) = \alpha V_t^{\text{IoT}} - \beta V_t^{\text{loss}} - \kappa V_t^{\text{below limit}}$$

- IoT traffic served
- Loss to conventional application after adding IoT traffic
- Bytes served below minimum throughput L

Implementation setup


- Deep Deterministic Policy Gradient (DDPG) algorithm [1]
 - Two hidden layers of sizes 400 and 300
 - Learning rate for actor and critic networks: 0.0001 and 0.001,
 - The discount factor is 0.99
 - Mini-batch size is 32
 - LSTM: two-layer architecture with 50 units per layer

[1] https://bitbucket.org/sandeep_chinchali/aaai18_deeprlcell

Utilization gain

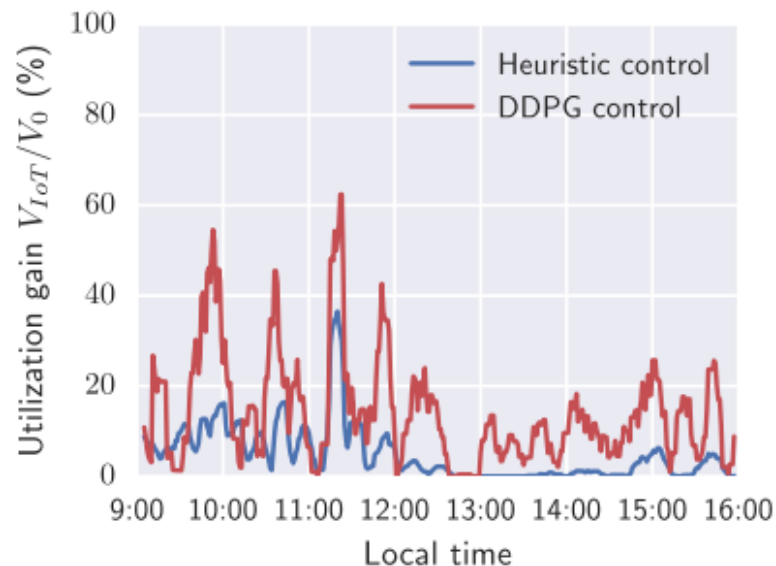
- Melbourne cell-day pairs (19 train, 8 test days)
 - IoT traffic favoring policy ($\alpha = 2, \beta = 1, \kappa = 1$)
 - Conservative policy ($\alpha = 1, \beta = 1, \kappa = 1$)

- Utilization gain: 14.7%




- The 10 MHz of radio spectrum, costing roughly \$4.5B,
 - Utilization gain of 14.7% means the operator is saved about **\$661 million**

RL vs. Benchmark Controllers



(a) IoT traffic



(b) Throughput

Thank you very much for your
attention.

Questions & Discussion...?